


Practical 06

Aim - Healthcare Analytics using logistic regression. To predict the likelihood of patient readmission based on key health indicator.

```
#importing libraries
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
#load the dataset
from google.colab import files
upload = files.upload()
```

 No file chosen Upload widget is only available when
Saving patient_data_large.csv to patient_data_large (1).csv

```
[ ] #converting cvs to dataframe
data = pd.read_csv('/content/patient_data_large.csv')
```

```

▶ print(data.info())#check for null values and data type
print(data.describe())#Summary statistics

```

```

⇒ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 7 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   age                                    1000 non-null   int64
1   gender                                1000 non-null   object
2   diabetes                              1000 non-null   int64
3   hypertension                          1000 non-null   int64
4   previous_admissions                  1000 non-null   int64
5   length_of_stay                       1000 non-null   int64
6   readmission                          1000 non-null   int64
dtypes: int64(6), object(1)
memory usage: 54.8+ KB
None

```

	age	diabetes	hypertension	previous_admissions \
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	52.881000	0.51000	0.502000	4.407000
std	20.958915	0.50015	0.500246	2.877087
min	18.000000	0.00000	0.000000	0.000000
25%	34.750000	0.00000	0.000000	2.000000
50%	52.500000	1.00000	1.000000	4.000000
75%	71.000000	1.00000	1.000000	7.000000
max	89.000000	1.00000	1.000000	9.000000

	length_of_stay	readmission
count	1000.000000	1000.000000
mean	7.652000	0.507000
std	4.046142	0.500201
min	1.000000	0.000000
25%	4.000000	0.000000
50%	8.000000	1.000000
75%	11.000000	1.000000
max	14.000000	1.000000

```
[ ] features = ['age', 'gender', 'diabetes', 'hypertension', 'previous_admissions', 'length_of_stay']  
target = 'readmission' #binary outcome: 1 for readmitted, 0 for not
```

```
X = data[features]  
y = data[target]
```

```
[ ] X = pd.get_dummies(X, drop_first=True) #convert categorical variables
```

```
[ ] scaler = StandardScaler()  
X_scaled = scaler.fit_transform(X)
```

```
[ ] X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.3, random_state=42)
```

```
[ ] model= LogisticRegression()  
model.fit(X_train, y_train)
```



```
▼ LogisticRegression  
LogisticRegression()
```

```
[ ] model_rf = RandomForestClassifier(n_estimators=100, random_state=42)  
model_rf.fit(X_train, y_train)
```



```
▼ RandomForestClassifier  
RandomForestClassifier(random_state=42)
```

```
[ ] y_pred = model.predict(X_test)
```

```

▶ accuracy = accuracy_score(y_test, y_pred)
  print("Accuracy:", accuracy)

  cm = confusion_matrix(y_test, y_pred)
  print("Confusion Matrix:\n", cm)

  report = classification_report(y_test, y_pred)
  print("Classification Report:\n", report)

```

```

⇒ Accuracy: 0.5266666666666666
Confusion Matrix:
[[67 87]
 [55 91]]
Classification Report:

```

	precision	recall	f1-score	support
0	0.55	0.44	0.49	154
1	0.51	0.62	0.56	146
accuracy			0.53	300
macro avg	0.53	0.53	0.52	300
weighted avg	0.53	0.53	0.52	300

```
[ ] y_pred_rf = model_rf.predict(X_test)
    accuract_rf = accuracy_score(y_test, y_pred_rf)
```

```

▶ from sklearn.model_selection import GridSearchCV

  param_grid = {'C': [0.1, 1, 10, 100], 'solver': ['lbfgs', 'liblinear']}
  grid = GridSearchCV(LogisticRegression(), param_grid, refit=True, verbose=3)
  grid.fit(X_train, y_train)

  print(grid.best_params_)

```

```

▶ Fitting 5 folds for each of 8 candidates, totalling 40 fits
[CV 1/5] END .....C=0.1, solver=lbfgs;; score=0.500 total time= 0.0s
[CV 2/5] END .....C=0.1, solver=lbfgs;; score=0.536 total time= 0.0s
[CV 3/5] END .....C=0.1, solver=lbfgs;; score=0.507 total time= 0.0s
[CV 4/5] END .....C=0.1, solver=lbfgs;; score=0.543 total time= 0.0s
[CV 5/5] END .....C=0.1, solver=lbfgs;; score=0.493 total time= 0.0s
[CV 1/5] END .....C=0.1, solver=liblinear;; score=0.493 total time= 0.0s
[CV 2/5] END .....C=0.1, solver=liblinear;; score=0.529 total time= 0.0s
[CV 3/5] END .....C=0.1, solver=liblinear;; score=0.514 total time= 0.0s
[CV 4/5] END .....C=0.1, solver=liblinear;; score=0.543 total time= 0.0s
[CV 5/5] END .....C=0.1, solver=liblinear;; score=0.493 total time= 0.0s
[CV 1/5] END .....C=1, solver=lbfgs;; score=0.493 total time= 0.0s
[CV 2/5] END .....C=1, solver=lbfgs;; score=0.529 total time= 0.0s
[CV 3/5] END .....C=1, solver=lbfgs;; score=0.507 total time= 0.0s
[CV 4/5] END .....C=1, solver=lbfgs;; score=0.543 total time= 0.0s
[CV 5/5] END .....C=1, solver=lbfgs;; score=0.493 total time= 0.0s
[CV 1/5] END .....C=1, solver=liblinear;; score=0.493 total time= 0.0s
[CV 2/5] END .....C=1, solver=liblinear;; score=0.529 total time= 0.0s
[CV 3/5] END .....C=1, solver=liblinear;; score=0.507 total time= 0.0s
[CV 4/5] END .....C=1, solver=liblinear;; score=0.543 total time= 0.0s
[CV 5/5] END .....C=1, solver=liblinear;; score=0.493 total time= 0.0s
[CV 1/5] END .....C=10, solver=lbfgs;; score=0.493 total time= 0.0s
[CV 2/5] END .....C=10, solver=lbfgs;; score=0.529 total time= 0.0s
[CV 3/5] END .....C=10, solver=lbfgs;; score=0.507 total time= 0.0s
[CV 4/5] END .....C=10, solver=lbfgs;; score=0.543 total time= 0.0s
[CV 5/5] END .....C=10, solver=lbfgs;; score=0.493 total time= 0.0s
[CV 1/5] END .....C=10, solver=liblinear;; score=0.493 total time= 0.0s
[CV 2/5] END .....C=10, solver=liblinear;; score=0.529 total time= 0.0s
[CV 3/5] END .....C=10, solver=liblinear;; score=0.507 total time= 0.0s
[CV 4/5] END .....C=10, solver=liblinear;; score=0.543 total time= 0.0s
[CV 5/5] END .....C=10, solver=liblinear;; score=0.493 total time= 0.0s

```

```

▶ [CV 4/5] END .....C=10, solver=liblinear;; score=0.543 total time= 0.0s
[CV 5/5] END .....C=10, solver=liblinear;; score=0.493 total time= 0.0s
[CV 1/5] END .....C=100, solver=lbfgs;; score=0.493 total time= 0.0s
[CV 2/5] END .....C=100, solver=lbfgs;; score=0.529 total time= 0.0s
[CV 3/5] END .....C=100, solver=lbfgs;; score=0.507 total time= 0.0s
[CV 4/5] END .....C=100, solver=lbfgs;; score=0.543 total time= 0.0s
[CV 5/5] END .....C=100, solver=lbfgs;; score=0.493 total time= 0.0s
[CV 1/5] END .....C=100, solver=liblinear;; score=0.493 total time= 0.0s
[CV 2/5] END .....C=100, solver=liblinear;; score=0.529 total time= 0.0s
[CV 3/5] END .....C=100, solver=liblinear;; score=0.507 total time= 0.0s
[CV 4/5] END .....C=100, solver=liblinear;; score=0.543 total time= 0.0s
[CV 5/5] END .....C=100, solver=liblinear;; score=0.493 total time= 0.0s
{'C': 0.1, 'solver': 'lbfgs'}

```

```

[] importances = model_rf.feature_importances_
feature_importance_df = pd.DataFrame({'Features': features, 'Importance': importances})
print(feature_importance_df.sort_values(by='Importance', ascending=False))

```

```

[]
      Features  Importance
0         age    0.404285
4  previous_admissions  0.245417
3       hypertension    0.211472
1         gender    0.047952
2        diabetes    0.045592
5    length_of_stay    0.045282

```